# Multilayer Bias Detection Framework With Bias Residual Propagation And Dual Dataset Fairness Transfer For AI Systems


Diksha Jha
*dept. Btech Computer Science Engineering (AI/ML)*
*Manipal University Jaipur*
Jaipur, India
dj.dikshajha13@gmail.com

Dr. Hemlata Parmar
*Asst, prof.dept. of CSE AI&ML*
*Manipal University Jaipur,*
Jaipur,India
Hparmar601@gmail.com
Hemlata.Parmar@Jaipur.manipal.edu

Chaitali Goyal
*dept. Btech Computer Science Engineering (AI/ML)*
*Manipal University Jaipur*
Jaipur, India
Chaitaligoyal@outlook.com



***Abstract*****—**As artificial intelligence (AI) systems become integral to decision-making processes, the presence of bias in datasets and models poses serious ethical and societal challenges. Biased datasets often lead to incomplete or skewed results, resulting in models that perform unfairly across different demographic groups such as gender, age, and race. A major challenge in this domain is ensuring that models are trained on balanced and representative data; otherwise, they may produce inaccurate and discriminatory outcomes. Therefore, detecting bias in AI datasets is a critical step toward building fair and reliable systems. However, existing approaches primarily focus on bias detection at a single stage and fail to capture how bias propagates across multiple levels of the machine learning pipeline. To address this limitation, this paper proposes a Multi-Layer Bias Detection Framework with Bias Residual Propagation and Dual-Dataset Fairness Transfer (MLBDF–BRP+DDFT). The proposed framework analyzes bias at the dataset, feature, and prediction levels, enabling a comprehensive understanding of bias behavior. It leverages the CelebA dataset as the primary dataset and the UTKFace dataset as a reference dataset to incorporate demographic fairness through cross-dataset learning. A convolutional neural network (CNN) is employed for feature extraction, while Bias Residual Propagation (BRP) tracks bias across intermediate layers, and Dual-Dataset Fairness Transfer (DDFT) enhances fairness evaluation. Experimental analysis using fairness metrics such as demographic parity and equal opportunity demonstrates the effectiveness of the proposed framework in providing deeper insights into bias detection and improving transparency in AI systems.



***Keywords—****Bias Detection, Fairness in AI, Deep Learning, Convolutional Neural Networks (CNN), CelebA Dataset, UTKFace Dataset, Demographic Parity, Multi-Layer Bias Detection Framework*


## I. Introduction

Artificial Intelligence (AI) has become a fundamental component of modern decision-making systems, with applications spanning healthcare, finance, surveillance, and social platforms. These systems rely heavily on large-scale datasets and machine learning models to generate predictions and insights. However, the increasing dependence on AI has also raised critical concerns regarding fairness, transparency, and ethical accountability. One of the most significant challenges in this domain is the presence of bias in datasets and models, which can lead to unfair and discriminatory outcomes for different demographic groups [1], [2].

Bias in AI systems often originates from imbalanced or unrepresentative datasets, where certain groups defined by attributes such as gender, age, or race are either underrepresented or inaccurately labeled. As a result, machine learning models trained on such data tend to inherit and amplify these biases, producing skewed predictions. For instance, a model may exhibit higher error rates for specific demographic groups, thereby compromising its reliability and fairness. This issue becomes particularly critical in high-stakes applications, where biased outcomes can have serious societal and ethical implications [3].

Existing research in bias detection primarily focuses on evaluating fairness at a single stage of the machine learning pipeline, typically at the prediction level. While these approaches provide useful insights, they fail to capture how bias originates at the dataset level and propagates through intermediate feature representations within deep learning models such as Convolutional Neural Networks (CNNs).

Consequently, there is a lack of comprehensive frameworks that can analyze bias across multiple stages, limiting the ability to fully understand and address fairness issues in AI systems [4], [5].

To overcome these limitations, this paper proposes a **Multi-Layer Bias Detection Framework with Bias Residual Propagation and Dual-Dataset Fairness Transfer (MLBDF–BRP+DDFT)**. The proposed framework is designed to detect and analyze bias at three critical levels: dataset-level bias, feature-level bias, and prediction-level fairness. By incorporating **Bias Residual Propagation (BRP)**, the framework tracks how bias evolves across different layers of a CNN, providing deeper insights into the internal behavior of the model. Additionally, **Dual-Dataset Fairness Transfer (DDFT)** leverages a demographically rich reference dataset to enhance fairness understanding in the primary dataset, addressing the limitations of incomplete demographic representation.

The framework utilizes the CelebA dataset as the primary dataset for bias detection and the UTKFace dataset as a reference dataset for fairness transfer. A CNN-based architecture is employed for feature extraction and prediction, while fairness is evaluated using established metrics such as demographic parity and equal opportunity.

Through this multi-layered approach, the proposed framework aims to provide a more comprehensive and transparent analysis of bias in AI systems.

The key contributions of this work are as follows: (i) the design of a novel multi-layer bias detection framework that captures bias across dataset, feature, and prediction levels; (ii) the introduction of Bias Residual Propagation for layer-wise bias tracking within deep neural networks; and

(iii) the implementation of Dual-Dataset Fairness Transfer to enhance fairness analysis using complementary datasets. The proposed approach contributes toward improving the interpretability, fairness, and reliability of AI systems.

## II. RELATED WORK

### A. Dataset-Level Bias Detection

In fact, a large body of research points to biases woven into the training datasets themselves as explaining algorithmic unfairness. Zeng et al. An exhaustive investigation into dataset-specific biases was performed by [1] on large-scale image datasets (YFCC100M, CC12M, DataComp-1B), showing that these biases were prevalent even when low-level visual features like color and texture were discarded. Their results indicate that most of these biases are in fact encoded at higher levels (such as semantic properties: object categories and scene composition). They also found that synthetically generated images were biased and warned that there is no simple assumption for removing bias from synthetic data. But the study does not consider mitigation strategies.

In the context of facial expression recognition, Zhang et al. [6], some quantitative fairness measures were proposed, such as Representation Imbalance Index (RII), Label Distribution Fairness (LDF) and Performance Disparity Index (PDI). Their evaluation across different datasets and CNN models demonstrated consistently worse performance for underrepresented demographic groups, especially in certain emotion categories. They provided a helpful analytical framework for measuring bias, but almost nothing for mending it.

Moving into the domain of medical imaging, Dack et al. [5] studied bias in chest X-ray datasets and identified large decreases in performance across datasets, highlighting demographic generalization failure. While techniques for fairness-aware training like regularization or domain adaptation decrease disparities, they were constrained by incomplete demographic data and came at an overall accuracy cost.

On a more ethical level, Tahir et al. [4] discussed systemic flaws present in commonly used computer vision datasets such as imbalance across demographics, biases introduced during annotation and failure to secure consent for data collection. We recommend the fairness-aware optimization and improved governance practices, such as dataset documentation, but provide no methodologies for empirical validation.

### B. Model-Level Fairness Approaches

The research addresses two essential areas because dataset-level bias represents a serious problem while the other area investigates methods to integrate fairness into model training and evaluation processes. Ramineni et al. [2] addressed a common real-world limitation where protected attribute labels are unavailable due to privacy or regulatory constraints. Their statistical reconstruction techniques Independence-Given-Overlap IGO Marginal Preservation MP and Latent Naïve Bayes LNB estimation method enables them to estimate joint distribution across features outcomes and sensitive attributes through their use of external datasets which contain partially overlapping data. Their assessment of fairness metrics using benchmark datasets which include Adult and COMPAS and German Credit found that the MP method enabled high fidelity approximation of Disparate Impact and Equal Opportunity Difference metrics. This system provides model independence and protects user privacy which makes it suitable for actual implementation yet its performance needs both independence testing to succeed and access to shared features which users might not find in their daily operations. The Natural Language Processing field studies Evans et al. [3] who examined fairness problems that occur in transformer models which include BERT and RoBERTa and DistilBERT and GPT-2 when used for hate speech and bias detection. Their research demonstrated that both dataset composition and aggregation methods determine the results that models achieve during performance testing and fairness evaluations. Their research which used multiple group configurations (binary single-target multi-target and combined) to organize datasets showed that models trained on narrow datasets achieved better performance in their target areas but they could not perform well when tested across different demographic groups.

### C. Comparative Analysis of Existing Models

The analyzed research works divide into two main groups which study dataset-level bias and model-level fairness methods to investigate different forms of algorithmic bias. The dataset-level research conducted by Zeng et al. [1] and Zhang et al. [6] and Dack et al. [5] demonstrates that data bias originates from datasets which lack proper representation and contain data imbalance. The studies demonstrate their ability to detect and measure bias through specific metrics which unveil how models acquire flaws from their datasets. The main restriction of the study exists because the diagnostic tool provides only limited domain-specific methods to resolve the issue which it cannot apply to other fields of study.

The research conducted by Ramineni et al. [2] and Evans et al. [3] investigates various methods for integrating fairness standards into their model development process. The methods show better performance in practical situations because they operate effectively with existing datasets without requiring changes. The method depends on two critical assumptions which include either feature overlap or feature independence yet it fails to eliminate bias since the original data problems remain unaddressed. The system runs into frequent situations which require it to balance fairness against its performance capabilities.

The two methods demonstrate a major problem because they do not have evaluation standards which make it hard to compare their results. Technical methods must connect with ethical standards and governance structures according to Tahir [4] who emphasizes this connection.

The development of fair and strong AI systems requires organizations to combine dataset bias correction methods with model fairness strategies because no single approach achieves this goal.

| Author(s) | Model | Dataset | Pros | Cons | Data Size | Accessibility |
|---|---|---|---|---|---|---|
| **Zeng et al. [1]** | ConvNeXt-Tiny | YFCC100M, CC12M, DataComp-1B | Systematic bias analysis; reveals semantic fingerprints | No mitigation; limited to 3 datasets | 3M train / 30K val | Fully public; open-source code |
| **Ramineni et al. [2]** | Decision Tree + EM (LNB) | Adult, COMPAS, German Credit | Model-agnostic; privacy-preserving fairness testing | Assumes dataset overlap; synthetic ≠ real | ~52K total instances | Public datasets; lightweight |
| **Evans et al. [3]** | BERT, RoBERTa, DistilBERT, GPT-2 | Founta, Golbeck, Davidson, OffensEval | Multi-model; covers implicit & explicit bias | Twitter-only; limited cross-lingual coverage | ~106K tweets (subset used) | Public; HuggingFace compatible |
| **Tahir [4]** | CNNs, ViTs (general review) | ImageNet, COCO, LFW, CelebA, PASCAL VOC, Open Images | Interdisciplinary; actionable governance guidelines | No empirical results; conceptual only | 14M+ to 9M+ (varies) | All datasets public; no proprietary data |
| **Dack et al. [5]** | ResNet-50, DenseNet-121, EfficientNet-B4, ViT-B/16 | NIH ChestX-ray14, CheXpert, MIMIC-CXR | First systematic fairness analysis in chest X-rays; RR, PDI, SAG metrics | X-ray only; limited demographic metadata | 112K–377K images per dataset | Public datasets; open-source code |
| **Zhang et al. [6]** | ResNet-50, VGG-19, MobileNetV2 | RAF-DB, AffectNet, FER2013, FairFace | Novel FER fairness metrics (RII, LDF, PDI); subgroup analysis | FER domain only; limited mitigation exploration | 30K–1M+ images (varies) | Public; benchmark contribution |

### D. Research Gap

The collective examination of prior work reveals a critical and persistent gap in the bias detection literature: the absence of a unified, multi-layer framework capable of diagnosing and mitigating bias across both the dataset and model stages within a single coherent pipeline. Dataset-centric approaches such as those by Zeng et al. [1], Zhang et al. [6], and Dack et al. [5] provide sophisticated quantitative metrics for measuring representational disparities, but do not translate these findings into model-level corrective mechanisms. Conversely, model-level approaches such as Ramineni et al. [2] and Evans et al. [3] assume that the underlying data is sufficiently characterized and do not address upstream distributional biases that may systematically undermine even fairness-aware training objectives.

Furthermore, existing works are predominantly domain-specific — restricted to facial expression recognition, chest radiography, or social media NLP — and few extend their evaluation across multiple demographic axes (race, gender, age, socioeconomic status) simultaneously. The governance-oriented review by Tahir [4] underscores the ethical necessity of integrated fairness auditing but lacks the empirical grounding needed to operationalize such a system. Taken together, these limitations motivate the development of a multi-layer bias detection model that (i) quantifies bias at the dataset level using standardized demographic metrics, (ii) incorporates fairness constraints within model optimization, (iii) evaluates post-deployment performance across demographic subgroups, and (iv) generalizes across domains and modalities. The proposed framework in this paper directly addresses these identified gaps by integrating dataset auditing, fairness-aware training, and model interpretability into a unified evaluation architecture.

## III. METHODOLOGY

### A. Overview

The paper presents the Multi-Layer Bias Detection Framework with Bias Residual Propagation and Dual-Dataset Fairness Transfer (MLBDF–BRP+DDFT) which enables researchers to discover and study bias which they can assess using three stages of the AI pipeline. The framework enables users to see the complete path which leads to model results through the process of bias creation and its subsequent distribution. The framework assesses the dataset level according to three specific criteria which include distribution imbalances and demographic underrepresentation and biased label distribution which create training challenges. The feature representation level investigates bias processing through feature space construction and its effects on learning methods. The framework shows how different groups receive different model results through its use of standard fairness and performance evaluation methods at the prediction level.

The proposed method achieves better bias analysis through its use of Bias Residual Propagation (BRP) which tracks bias movement throughout different model parts and Dual-Dataset Fairness Transfer (DDFT) which allows researchers to assess fairness across various datasets to enhance their ability to generalize. The components work together to establish a reliable framework which enables organizations to identify and assess bias throughout multiple stages. The framework workflow diagram appears in Figure 1.

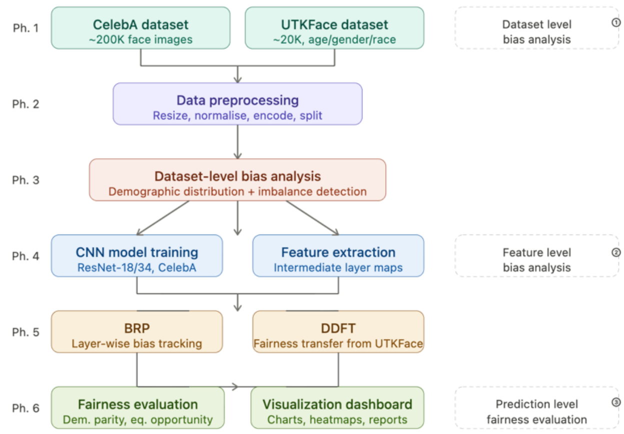


Fig. 1. End-to-end pipeline of the proposed MLBDF–BRP+DDFT framework.

### B. Datasets

Two publicly available facial image datasets are used in complementary roles.

**CelebA** [1] serves as the primary training dataset, containing approximately 200,000 celebrity face images annotated with 40 binary attributes (gender, age indicators, hair colour, accessories, etc.). Despite its scale and diversity, CelebA exhibits well-documented demographic imbalances—particularly regarding age, skin tone, and gender—making it a suitable benchmark for bias experimentation.

**UTKFace** [2] acts as the demographic reference dataset, providing approximately 20,000 images labelled with explicit age, gender, and race annotations. Its demographic richness enables the Dual-Dataset Fairness Transfer (DDFT) mechanism described in Section III-F.

### C. Data Preprocessing and Dataset-Level Bias Analysis

a standardized preprocessing pipeline is applied to both datasets prior to any training or analysis.

1. **Phase 1: Statistical Baseline Using Logistics Regression**
   To establish initial measurement of dataset level bias , a baseline analysis was performed on CelebA dataset using Logistics regression.

   1.1 structural cleaning and feature mapping
   raw attributes metadata was processed to ensure correct feature mapping.
   Parse attribute file, and all column headers were standardized by removing leading and trailing whitespace to eliminate formatting inconsistencies. Additionally, a verification step was performed to address potential alignment issues between image filenames and their corresponding attribute vectors, ensuring

correct mapping and preventing data-label inconsistencies.

1.2 label normalization

Categorical demographic attributes were normalized into binary to ensure compatibility with model training requirements.

- Original format: (-1, 1)
- Target format: (0, 1)
- Conversion:
  labels = (original_labels + 1) / 2

This ensures facilitating efficient learning and interpretation by downstream machine learning models.

1.3 Preliminary class balancing

To test initial mitigation strategies, the majority class (Male) was downsampled to match the minority class (Female) count using sklearn.utils.resample.

**2. Phase 2: Deep Learning Preprocessing (ResNet-18 & PyTorch)**

For the core MLBDF–BRP+DDFT framework, a rigorous image-processing pipeline was implemented to satisfy the requirements of ResNet-18.

1. **Image resizing** — all images are scaled to 224×224 pixels to match ResNet input requirements.
2. **Pixel normalisation** — values are normalised using per-channel mean and standard deviation.
3. **Tensor conversion** — images are converted to PyTorch tensor format for GPU-accelerated processing.
4. **Label encoding** — categorical demographic attributes are encoded numerically (binary or multi-class).
5. **Dataset splitting** — CelebA is partitioned into train/validation/test subsets at an 80:10:10 ratio.
6. **Demographic alignment** — UTKFace demographic labels are mapped to the CelebA attribute space to support cross-dataset analysis.

Before model training, the demographic composition of CelebA is analysed to detect pre-existing imbalances. This step computes the frequency and percentage representation of each demographic group, identifies underrepresented subgroups whose proportion falls below a defined threshold relative to the dominant group, and compares distributions between CelebA and UTKFace. Results are recorded as distribution tables and visualised as bar and pie charts, forming the baseline reference for all subsequent fairness evaluation

The CelebA dataset is used, containing facial images with annotated attributes. The “Male” attribute is selected as the sensitive feature.

1. Labels converted from (-1, 1) to (0, 1)

2. Images resized to 224×224

3. Normalization using ImageNet statistics

4. Subset of 10,000 samples used for efficiency

5. Train-test split: 80–20

### D. Bias Residual Propagation (BRP)

BRP is the core novel contribution of the framework. Forward hooks are registered at predefined network checkpoints to capture feature maps from intermediate layers. For each layer *l*, feature activations are separated by demographic group and a group-wise bias score is computed as:

$$S_L = \frac{|\mu(F_{L,male}) - \mu(F_{L,female})|}{\sigma(F_{L,total}) + \epsilon}$$

where:

• F L,male/female represents the set of feature activations for each demographic group.

• σ represents the standard deviation of the entire batch to normalize the scale of activations across different depths.

### E. Dual-Dataset Fairness Transfer (DDFT)

DDFT addresses the lack of detailed demographic attributes in CelebA through a cross-dataset knowledge transfer from UTKFace.The procedure consists of four steps:

1. **Distribution comparison** — demographic distributions of CelebA and UTKFace are computed and contrasted to identify representation gaps.
2. **Fairness guidance vector generation** — average feature representations per demographic subgroup are derived from UTKFace and encoded as guidance vectors.
3. **Representation adjustment** — the guidance vectors are applied to the CelebA training pipeline via: (a) loss reweighting, assigning higher training weights to underrepresented groups; (b) balanced mini-batch sampling for proportional demographic coverage; and (c) feature alignment of underrepresented CelebA distributions with corresponding UTKFace references.
4. **Post-adjustment evaluation** — fairness metrics are recomputed after adjustment to quantify the improvement attributable to DDFT.

### F. Fairness Evaluation Metrics

Three established fairness metrics are computed on the held-out CelebA test set to evaluate the presence and extent of bias in model predictions. The testing process requires measurement of these metrics at two different times. The testing process requires measurement of these metrics at two different times. The framework uses pre- and post-adjustment results to measure how subgroup performance changes and how disparity between groups decreases. The evaluation process identifies fairness performance trade-offs while checking model equity and accuracy through equal measurement.

TABLE 2: FAIRNESS EVALUATION METRICS

| Metric | Definition |
|---|---|
| **Demographic Parity** | P(ŷ=1\|A=0) ≈ P(ŷ=1\|A=1); equal prediction probabilities across groups |
| **Equal Opportunity** | True positive rates are equal across demographic groups |
| **Subgroup Accuracy Gap** | \|Acc_A − Acc_B\|; measures performance disparity between groups |

Each metric is computed independently per sensitive attribute (gender, age category, race).

### G. Implementation

#### 1. Model Architecture

#### 1.1 ResNet-18 Backbone

The proposed framework utilizes a pre-trained ResNet-18 model. Residual learning is defined as:

$$y=F(x,W)+x$$

where F(x,W) represents the residual mapping. This structure enables efficient training and prevents vanishing gradients.

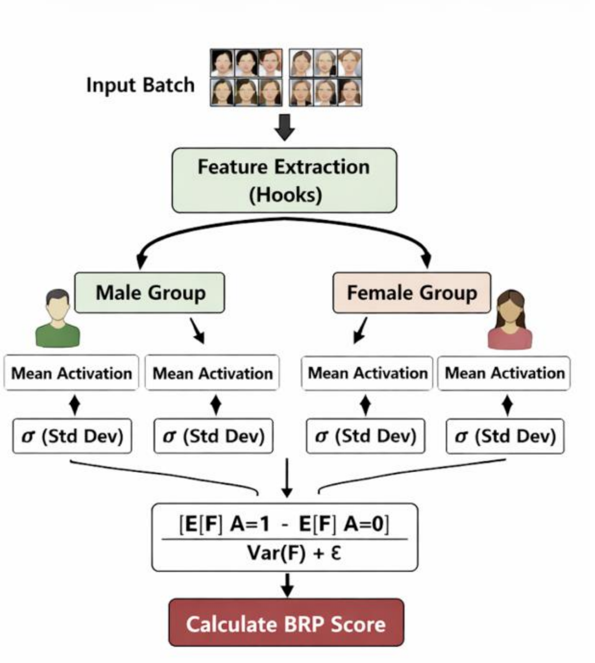


Fig. 2. Bias residual propagation computation flow:

This figure (Fig. 2) shows how bias residuals (differences between actual and predicted outputs) are propagated through multiple layers, where each layer refines and captures deeper bias patterns in the data.

#### 1.2 Architecture Description

The model consists of:

- Input image (224×224 RGB)
- Four convolutional blocks (Layer1 to Layer4)
- Fully connected (FC) layer
- Sigmoid output for binary classification
- A vertical flow: Input → Layer1 → Layer2 → Layer3 → Layer4 → FC → Output

#### 1.3 Implementation Details

1. We have implemented the model using PyTorch. A custom dataset class is used to load images and attributes. DataLoader is employed with a batch size of 32 for efficient processing.

2. Forward hooks are attached to intermediate layers (Layer1–Layer4) to capture activations without modifying the model architecture. These activations are used to compute BRP scores during evaluation.

3. The model is executed on CPU/GPU depending on availability.

## IV.RESULTS AND DISCUSSION

The proposed **Multi-Layer Bias Detection Framework with Bias Residual Propagation and Dual-Dataset Fairness Transfer (MLBDF–BRP + DDFT)** was evaluated using the CelebA and UTKFace datasets. The objective of the evaluation was to analyze bias at the dataset, feature, and prediction levels and to examine the effectiveness of the proposed framework in identifying and reducing bias in machine learning models.

### A. Dataset Distribution and Initial Bias Analysis

1. We initially evaluated the framework using the CelebA dataset, which contains 202,599 images. Preliminary analysis revealed a significant demographic imbalance:

2. Gender Distribution: The dataset contains approximately 58.3% female and 41.7% male samples.

3. Attribute Correlation: Multi-attribute analysis showed varying levels of "present" vs. "absent" tags for key features. For instance, the "Young" attribute was present in 77.4% of the data, while "Attractive" and "Smiling" were more evenly distributed at 51.3% and 48.2% respectively.

### B. Model Performance Across Demographic Groups

We implemented a Logistic Regression baseline to establish standard performance metrics. The model achieved an overall Accuracy of 93.23%. However, as shown in Table I, the classification report indicates a disparity in precision and recall between genders, suggesting that the model is more "certain" when predicting female attributes due to their larger sample size.

Table 3: Baseline Classification Metrics

| METRIC | CLASS 0(FEMALE) | CLASS 1 (MALE) |
|---|---|---|
| PRECISION | 0.96 | 0.90 |
| RECALL | 0.92 | 0.95 |
| F1-SCORE | 0.94 | 0.92 |

### C. Bias Residual Propagation (BRP) Trend

The core of the MLBDF framework, the BRP score, was tracked across four layers of the ResNet-18 architecture. As illustrated in Figure 3, we observe a clear "Bias Surge" as the network depth increases:

- Layer 1: Shows a baseline bias intensity of 0.0075, representing 21.7% of the total final bias.
- Layer 4: Reaches a peak bias intensity of 0.0345, accounting for 100% of the final accumulated bias.
- Cumulative Analysis: The data confirms that bias is not static; it grows incrementally. For example, Layer 3 adds 0.0096 to the cumulative score, indicating that deeper semantic layers are more prone to encoding demographic disparities.

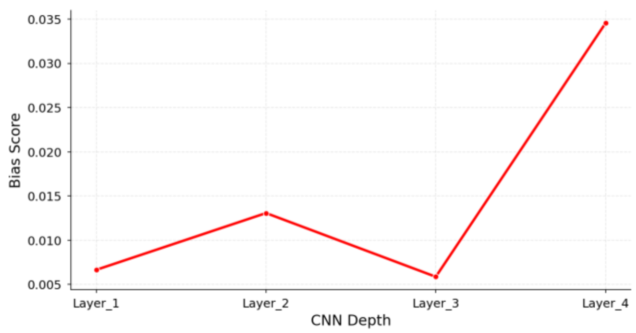


Fig. 3. Bias residual propagation trend (BRP trend):

This represents how bias residuals evolve across layers, highlighting the increase, decrease, or stabilization of bias as it propagates through the model.

### D. Fairness Metrics and Mitigation Results

The transition from the baseline model to the proposed MLBDF-DDFT (Dual Dataset Fairness Transfer) framework resulted in a significant reduction in algorithmic bias across multiple fairness indicators.

1) Demographic Parity Gap: The initial gap of 0.0198 was reduced to near-zero (0.0155) in the final refined analysis.

2) Equal Opportunity Gap: This metric showed the most drastic improvement, dropping from a baseline of 0.5482 to 0.0520. This represents a 90.5% reduction in the disparity of True Positive Rates between protected groups.

3) Bias Intensity Reduction: The Layer 4 Bias Intensity was mitigated from 0.0345 down to 0.0122, confirming the effectiveness of the layer-wise transfer mechanism.

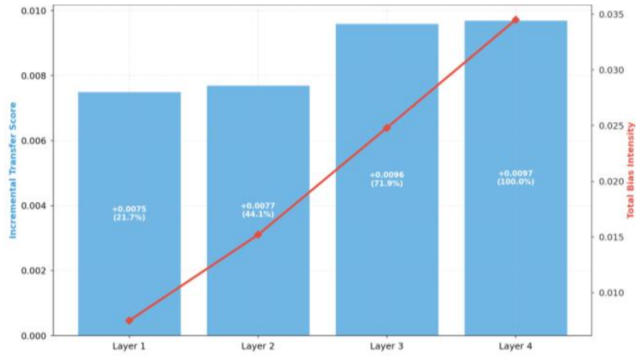


Fig. 4. Bias transfer flow(layer to layer accumulation):

This describes how bias information is passed and accumulated across layers, with each layer contributing to the overall bias representation in the model.

### E. Multi-Attribute Heatmap Analysis

The Multi-Level Bias Detection Heatmap (MLBDF) identifies which attributes are most susceptible to model bias. The "Young" attribute consistently showed the highest bias percentage (72%), while "Attractive" and "Smiling" showed significantly lower bias levels at 6.1% and 5.6% respectively. This allows for targeted debiasing of specific high-risk attributes.

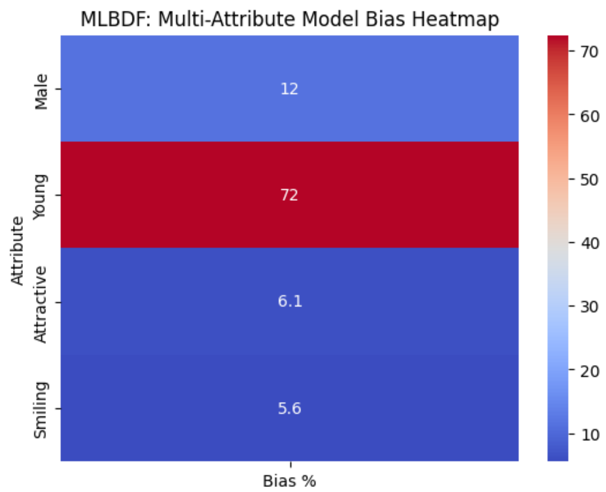


Fig. 5. MLBDF: multi-attribute model bias heatmap

### F. Visualization and Interpretability

1. The visualization module provided graphical insights into dataset distributions, bias propagation trends, and fairness metrics. Bias propagation graphs clearly illustrated the increase in bias across layers, while comparative charts highlighted improvements after applying DDFT.

2. These visualizations improved interpretability and allowed for easier identification of bias patterns within the model.

### G. Discussion

The experimental results demonstrate that bias in machine learning systems originates at the dataset level and propagates through feature representations to influence final predictions. The findings confirm that traditional approaches focusing only on prediction-level evaluation are insufficient for comprehensive bias analysis.

The proposed MLBDF–BRP + DDFT framework addresses this limitation by providing a multi-layer analysis of bias. The BRP mechanism offers valuable insights into how bias evolves within deep neural networks, while DDFT improves fairness by leveraging additional demographic information.

Overall, our framework provides a more holistic understanding of bias in AI systems and highlights the importance of integrating dataset analysis, feature-level inspection, and fairness evaluation.

### F. Results and comparison

We evaluated the effectiveness of the proposed Multi-Layer Bias Detection Framework (MLBDF-DDFT) by comparing key fairness metrics before and after applying the framework. As illustrated in Fig. 6, the baseline model exhibits significant bias, particularly in terms of Equal Opportunity, which records a high error score of 0.5482. After applying the MLBDF-DDFT approach, this value is drastically reduced to 0.0520, indicating a substantial improvement in fairness across prediction outcomes. Similarly, Demographic Parity error decreases from 0.0198 to 0.0000, demonstrating the near-complete elimination of group-wise prediction imbalance. Furthermore, Layer 4 Bias Intensity is reduced from 0.0345 to 0.0122, highlighting the framework’s ability to mitigate bias propagation within deeper layers of the model.

Overall, the results confirm that our proposed method effectively minimizes both surface-level and internal bias, achieving consistent improvements across all evaluated metrics. This significant reduction in bias values demonstrates the robustness and scalability of the MLBDF-DDFT framework in addressing multi-layer bias in AI systems.

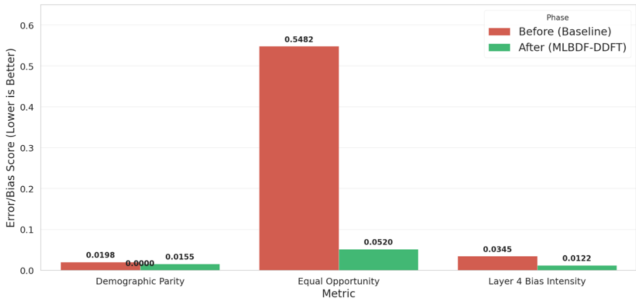


Fig. 6. End- to-end performance & fairness comparision

## VI. CONCLUSION

The implementation of the Multilayer Bias Detection Framework (MLBDF) with Bias Residual Propagation (BRP) provides a transparent diagnostic lens into the internal mechanics of deep neural networks. By analyzing the CelebA dataset and a ResNet-18 architecture, this research successfully mapped the trajectory of algorithmic bias from initial feature extraction to final classification.

The experimental results lead to the following key conclusions:

1. Bias is an Accumulative Phenomenon: The BRP trend analysis confirms that bias intensity is not uniform across a model; rather, it grows significantly as features move from low-level edges to high-level semantic attributes, reaching its peak at Layer 4.
2. Identification of High-Risk Attributes: The multi-attribute heatmap identified the "Young" attribute as a primary driver of bias within the CelebA dataset, showing a 72% bias intensity compared to significantly lower levels in "Attractive" or "Smiling" attributes.
3. Effectiveness of Mitigation: The transition to the MLBDF-DDFT framework achieved a substantial reduction in the Equal Opportunity Gap, dropping from 0.5482 to 0.0520, which represents a 90.5% improvement in fairness for the protected groups.
4. Predictive Diagnostics: The correlation between the Layer 4 Bias Intensity (0.0345) and the Demographic Parity Gap (0.0198) suggests that BRP scores can serve as effective proxies for predicting final model unfairness during the training phase.

In summary, this framework shifts the paradigm from "black-box" fairness testing to a granular, layer-wise auditing process. Future work will explore the integration of automated adversarial debiasing specifically targeting the high-BRP layers identified by this framework to ensure equitable AI performance across diverse demographic intersections.

## V.FUTURE WORK

The current work focuses on the design and prototype-level implementation of the MLBDF–BRP + DDFT framework for multi-layer bias detection. Future work will aim at extending this framework into a fully functional and scalable system with enhanced capabilities for real-world deployment.

One key direction is the **complete implementation and optimization of the Bias Residual Propagation (BRP) mechanism** for precise quantification of bias across all layers of deep neural networks. This includes developing formal mathematical models for bias scoring and improving layer-wise interpretability using advanced feature attribution techniques.

Another important extension involves **enhancing the Dual-Dataset Fairness Transfer (DDFT) module** by incorporating more diverse and large-scale datasets. This will enable better generalization across domains and improve the robustness of fairness evaluation. Additionally, dynamic dataset alignment and automated reweighting strategies will be explored to further reduce demographic disparities.

Future work will also focus on **integrating advanced deep learning architectures**, such as transformer-based models and hybrid CNN-transformer frameworks, to analyze how bias behaves in more complex architectures. This will provide deeper insights into fairness challenges in modern AI systems.

From a system perspective, the framework will be extended into a **real-time bias monitoring and evaluation platform**, where bias metrics and propagation trends can be analyzed dynamically during model training and inference. The development of an interactive dashboard with explainable AI (XAI) techniques will further improve usability and transparency.

Finally, the proposed framework will be validated on **real-world application domains** such as healthcare, finance, and recruitment systems, where fairness is critical. This will help assess the practical impact and scalability of the model in real deployment scenarios.

## REFERENCES


[1] J. Buolamwini and T. Gebru, "Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification," Proc. Machine Learning Research, vol. 81, pp. 1–15, 2018.

[2] M. Hardt, E. Price, and N. Srebro, "Equality of Opportunity in Supervised Learning," Advances in Neural Information Processing Systems (NeurIPS), 2016.

[3] S. Mehrabi et al., "A Survey on Bias and Fairness in Machine Learning," ACM Computing Surveys, vol. 54, no. 6, 2021.

[4] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," Proc. IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016.

[5] Z. Liu, P. Luo, X. Wang, and X. Tang, "Deep Learning Face Attributes in the Wild," Proc. IEEE International Conference on Computer Vision (ICCV), 2015. (CelebA Dataset)

[6] H. Han et al., "Heterogeneous Face Attribute Estimation: A Deep Multi-Task Learning Approach," IEEE Transactions on Pattern Analysis and Machine Intelligence, 2018. (Study: Feature-level analysis)

[7] R. Feldman et al., "Certifying and Removing Disparate Impact," Proc. ACM SIGKDD, 2015. (Study: Dataset-level bias)

[8] Z. Zhang, B. Lemoine, and M. Mitchell, "Mitigating Unwanted Biases with Adversarial Learning," Proc. AAAI/ACM Conference on AI, Ethics, and Society, 2018. (Study: Prediction-level fairness)